


Encoder (runs once per case)

Decoder (autoregressive, one step shown)



0.5–0.7 M parameters · vocab = 80 links + SEP/BOS/EOS/PAD · joint sequence $\approx V \times (\text{route len} + 1) + 2$ tokens · training: teacher forcing + cross-entropy on Gurobi reference sequences (mask not needed); inference: masked greedy / nucleus decode